

TMDB Box Office Prediction

Predicting the box office revenue for a given movie

Meet the Analysts



Carlos G. Fernandez



Luis Eduardo Leal



Nicole Ardizzi

Why this topic?

- Who doesn't love movies?
- \$41.7 billion industry in 2018,
- The film industry is more popular than ever.
- What if you can make a movie?
- Interesting features to analyze

The Dataset

Source	Description	Size
<p>Kaggle.com</p> <p>TMDB Box Office Prediction</p> <p>Can you predict a movie's worldwide box office revenue?</p>	<ul style="list-style-type: none">• A list of movies and a variety of metadata• Metadata from The Movie Database (TMDB)• Collected from the TMDB Open API	<ul style="list-style-type: none">• 3000 rows (movies)• 23 columns

Dataflow

Postgres → Pandas

- Raw data was stored in Postgres
- Transferred data from Postgres to Pandas using SQLAlchemy
- `pd.read_sql()`

Pandas → Tableau

- `.to_csv()`
-

Columns Names & Data Types

String

- Id
- Homepage
- Imbd_id
- Original language
- Original_title
- Overview
- Poster path
- Release_date
- Status
- Tagline
- Title

Numerical

- Budget
- Popularity
- Runtime
- Revenue

Data Structure Format

- Belongs to collection
- genres
- Production Companies
- Production_countries
- Spoken_languages
- Keywords
- Cast
- Crew



Exploratory Data Analysis / Data Processing

Exploratory Analysis

Questions

- How does budget influence revenue?
- What is the average revenue for each genre?
- Revenue of movies with home page vs no homepage?
- Did production companies have a major influence on revenue?
- Which production company had the highest average revenue?

Tools and Technologies Used During EDA

- Pandas
 - Numpy
 - Matplotlib
 - Seaborn
 - SQLAlchemy
 - Pearsonr
-

Genres

For loop function:

- Goes through each row
- Looks at each genre list
- Looks at each genre in the genre list

Unpack and encode:

- Unique genre names are converted into columns
- For each genre in the genre list, encode with 1 if column name = genre name.

```
# check new previous vs transformed column  
train.loc[:, ['genres', 'new_genre']]
```

	genres	new_genre
0	[{'id': 35, 'name': 'Comedy'}]	[Comedy]
1	[{'id': 35, 'name': 'Comedy'}, {'id': 18, 'name': 'Drama'}]	[Comedy, Drama, Family, Romance]
2	[{'id': 18, 'name': 'Drama'}]	[Drama]
3	[{'id': 53, 'name': 'Thriller'}, {'id': 18, 'name': 'Drama'}]	[Thriller, Drama]
4	[{'id': 28, 'name': 'Action'}, {'id': 53, 'name': 'Thriller'}]	[Action, Thriller]
...
2995	[{'id': 35, 'name': 'Comedy'}, {'id': 10749, 'name': 'Romance'}]	[Comedy, Romance]
2996	[{'id': 18, 'name': 'Drama'}, {'id': 10402, 'name': 'Music'}]	[Drama, Music]
2997	[{'id': 80, 'name': 'Crime'}, {'id': 28, 'name': 'Action'}, {'id': 53, 'name': 'Thriller'}]	[Crime, Action, Mystery, Thriller]
2998	[{'id': 35, 'name': 'Comedy'}, {'id': 10749, 'name': 'Romance'}]	[Comedy, Romance]
2999	[{'id': 53, 'name': 'Thriller'}, {'id': 28, 'name': 'Action'}, {'id': 80, 'name': 'Crime'}]	[Thriller, Action, Mystery]
3000 rows x 2 columns		

Cast

- JSON format
- Used regular expressions to clean data
- Created a dictionaries:
 - Actors dictionary
 - Lead actors (1st one credited for a film)
 - Top actors
 - Count of 'top actors' per movie. Top actors reference by top 50 actors in the cast

```
train.iloc[0]['cast']
```

✓ 0.1s

Python

```
{'cast_id': 4, 'character': 'Lou', 'credit_id': '52fe4ee7c3a36847f82afae7', 'gender': 2, 'id': 52997, 'name': 'Rob Corddry', 'order': 0, 'profile_path': '/kZzJL0V1nEZUF08xUd0d3ucfXz.jpg'}, {'cast_id': 5, 'character': 'Nick', 'credit_id': '52fe4ee7c3a36847f82afae7', 'gender': 2, 'id': 64342, 'name': 'Craig Robinson', 'order': 1, 'profile_path': '/tVaRMkJX0EVhYtnnFuhqW0Rjzz.jpg'}, {'cast_id': 6, 'character': 'Jacob', 'credit_id': '52fe4ee7c3a36847f82afae7', 'gender': 2, 'id': 54729, 'name': 'Clark Duke', 'order': 2, 'profile_path': '/oNzK0umwm5Wn0wyEb0y6TVJCSBn.jpg'}, {'cast_id': 7, 'character': 'Adam Jr.', 'credit_id': '52fe4ee7c3a36847f82afaf3', 'gender': 2, 'id': 36801, 'name': 'Adam Scott', 'order': 3, 'profile_path': '/5gb65xz8bzd42yjMAL4zwo4cvKw.jpg'}, {'cast_id': 8, 'character': 'Hot Tub Repairman', 'credit_id': '52fe4ee7c3a36847f82afaf7', 'gender': 2, 'id': 54812, 'name': 'Chevy Chase', 'order': 4, 'profile_path': '/svjpyYtPwtjvRxX9Izn0m0khD0t.jpg'}, {'cast_id': 9, 'character': 'Jill', 'credit_id': '52fe4ee7c3a36847f82afafb', 'gender': 1, 'id': 94098, 'name': 'Gillian Jacobs', 'order': 5, 'profile_path': '/rBnhe5vhNPNhRUdtYahBWx90fJM.jpg'}, {'cast_id': 10, 'character': 'Sophie', 'credit_id': '52fe4ee7c3a36847f82afaff', 'gender': 1, 'id': 1159009, 'name': 'Bianca Haase', 'order': 6, 'profile_path': '/4x3nbtD8q8phAJPmoGWXPv0iM.jpg'}, {'cast_id': 11, 'character': 'Kelly', 'credit_id': '5524ec51c3a3687df3000dbb', 'gender': 1, 'id': 86624, 'name': 'Collette Wolfe', 'order': 7, 'profile_path': '/aSD4h5379b2eW3bLou9ByLimmq.jpg'}, {'cast_id': 13, 'character': 'Brad', 'credit_id': '5524ec8ec3a3687ded000d72', 'gender': 2, 'id': 466505, 'name': 'Kumail Nanjiani', 'order': 9, 'profile_path': '/x4nAztHY72SVciRfxEsbhIVTSiu.jpg'}, {'cast_id': 14, 'character': 'Courtney', 'credit_id': '5524ec9bc3a3687df8000d13', 'gender': 1, 'id': 70776, 'name': 'Kellee Stewart', 'order': 10, 'profile_path': '/w3xmsEPmJc1Cf0dQ4aIn8YmLHbk.jpg'}, {'cast_id': 15, 'character': 'Terry', 'credit_id': '5524eca892514171cb008237', 'gender': 2, 'id': 347335, 'name': 'Josh Heald', 'order': 11, 'profile_path': '/pwXJ1IenrDMrG7t3zNFlvr8w1RGU.jpg'}, {'cast_id': 16, 'character': 'Susan', 'credit_id': '5524ecb7925141720c001116', 'gender': 0, 'id': 1451392, 'name': 'Gretchen Koerner', 'order': 12, 'profile_path': '/muULPexCTJGyJba4yKzxrnpD50.jpg'}, {'cast_id': 17, 'character': 'Herself', 'credit_id': '5524ecc3c3a3687ded000d74', 'gender': 1, 'id': 98879, 'name': 'Lisa Loeb', 'order': 13, 'profile_path': '/bGqg58ca0bZR38z9HliUmMeNGE.jpg'}, {'cast_id': 18, 'character': 'Herself', 'credit_id': '5524ecd3c3a3687e11000ed3', 'gender': 1, 'id': 1394648, 'name': 'Jessica Williams', 'order': 14, 'profile_path': '/A4syKjKcYB92wLEhH0c0hC3BCpz.jpg'}, {'cast_id': 19, 'character': 'Himself', 'credit_id': '5524ece6925141718d001009', 'gender': 1}
```

Crew


- JSON format
- Used regular expressions to clean data
- Created dictionaries:
 - Directors
 - Producers
 - Executive producers
 - Number of male/female crew members

```
train.iloc[1]['crew']  
✓ 0.6s Python  
"[{'credit_id': '52fe43fe9251416c7502563d', 'department': 'Directing', 'gender': 2, 'id': 1201, 'job': 'Director', 'name': 'Garry Marshall', 'profile_path': '/kx77E8p5rnEmKxIhFT0qWCEMEik.jpg'}, {'credit_id': '52fe43fe9251416c75025667', 'department': 'Camera', 'gender': 2, 'id': 1214, 'job': 'Director of Photography', 'name': 'Charles Minsky', 'profile_path': None}, {'credit_id': '52fe43fe9251416c75025661', 'department': 'Sound', 'gender': 2, 'id': 4500, 'job': 'Original Music Composer', 'name': 'John Debney', 'profile_path': '/hTrlvZLDXQk49nfc2BM9sjKfJv.jpg'}, {'credit_id': '52fe43fe9251416c7502564f', 'department': 'Production', 'gender': 1, 'id': 8851, 'job': 'Producer', 'name': 'Whitney Houston', 'profile_path': '/69ouDnXnmkLYPr4sMJXWKYz81AL.jpg'}, {'credit_id': '52fe43fe9251416c7502566d', 'department': 'Editing', 'gender': 0, 'id': 12970, 'job': 'Editor', 'name': 'Bruce Green', 'profile_path': '/yplxWPVgwK1b33AjvbmM9mwX2Aw.jpg'}, {'credit_id': '52fe43fe9251416c75025655', 'department': 'Production', 'gender': 2, 'id': 38415, 'job': 'Producer', 'name': 'Mario Iscovich', 'profile_path': None}, {'credit_id': '52fe43fe9251416c7502565b', 'department': 'Production', 'gender': 1, 'id': 38416, 'job': 'Executive Producer', 'name': 'Ellen H. Schwartz', 'profile_path': '/6WInjABr1sAYGXaa5q0vSrsHIqP.jpg'}, {'credit_id': '52fe43fe9251416c75025649', 'department': 'Production', 'gender': 1, 'id': 59973, 'job': 'Producer', 'name': 'Debra Martin Chase', 'profile_path': None}, {'credit_id': '52fe43fe9251416c75025643', 'department': 'Writing', 'gender': 1, 'id': 25539, 'job': 'Screenplay', 'name': 'Shonda Rhimes', 'profile_path': '/4c77e347InbTA1w9LgvORpZBHV6.jpg'}]"
```

EDA Questions

- **How does budget influence revenue?**
- What is the average revenue for each genre?
- Revenue of movies with home page vs no homepage?
- Did production companies have a major influence on revenue?
- Which production company had the highest average revenue?

Findings:

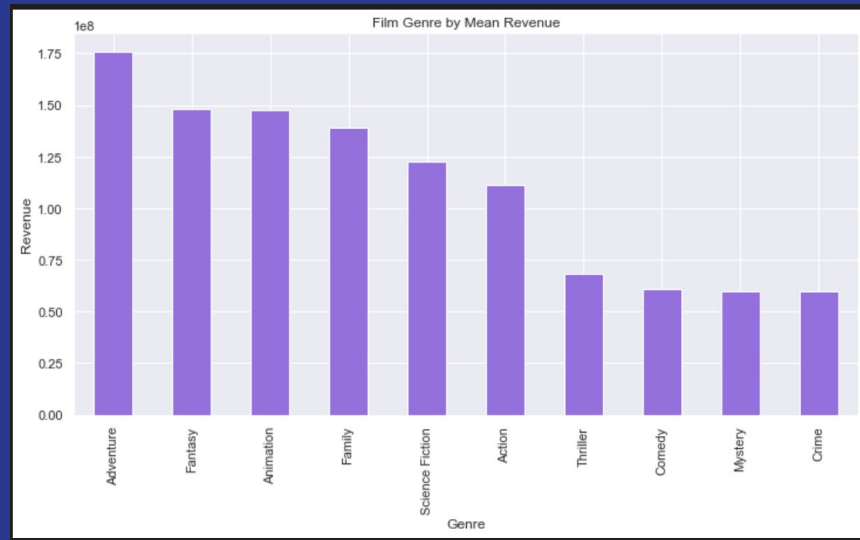


There is a strong correlation

EDA Questions

- How does budget influence revenue?
- **What is the average revenue for each genre?**
- Revenue of movies with home page vs no homepage?
- Did production companies have a major influence on revenue?
- Which production company had the highest average revenue?

Findings:



Adventure films have the highest revenue values

EDA Questions

- How does budget influence revenue?
- What is the average revenue for each genre?
- **Revenue of movies with home page vs no homepage?**
- Did production companies have a major influence on revenue?
- Which production company had the highest average revenue?

Findings:

Mean Revenue

Movies with a homepage:
\$120,051,698

No homepage:
\$42,165,846

EDA Questions

- How does budget influence revenue?
- What is the average revenue for each genre?
- Revenue of movies with home page vs no homepage?
- **Did production companies have a major influence on revenue?**
- Which production company had the highest average revenue?

Findings:

Based of the Pearson correlation:

- The number of top studios a movie uses in relation to revenue are strongly correlated.
 - $r = 0.558$


Note:

- Top studios typically have more reputation, funds and resources that benefits the movie overall
-

EDA Questions

- How does budget influence revenue?
- What is the average revenue for each genre?
- Revenue of movies with home page vs no homepage?
- Did production companies have a major influence on revenue?
- **Which production company had the highest average revenue?**

Findings:






Machine Learning Models

Selecting Features

Numerical Correlation Matrix

- Set 1: Categorical Inputs
 - Genre
 - Studio
 - Collection
 - Transformed into binary
- Set 2: Quantitative Inputs
 - Budget
 - Runtime
 - Popularity
- Set 3: Transforming Features
 - Log Transformation



Feature Importances

Target: Revenue

- Uses feature of importance to rank relevant metrics
- Some of the features that ranked top across the four different models are similar.

```
feature_importances = rfr_best_model.feature_importances_  
indices = np.argsort(feature_importances)[::-1]  
  
# Print the feature ranking  
print("Feature ranking:")  
  
for f in range(20):  
    print("%d. %s (%f)" % (f + 1, X_train.columns[indices[f]], feature_importances[indices[f]]))
```

✓ 0.4s

Feature ranking:

1. topStudio (0.361790)
2. usa_produced (0.117028)
3. belongs_to_collection (0.109838)
4. topLeadActor (0.097712)
5. has_homepage (0.060809)
6. fridayRelease (0.028559)
7. topLeadDirector (0.028427)
8. released_in_english (0.024508)
9. originally_english (0.020992)
10. Fall (0.018236)
11. Winter (0.018028)
12. wednesdayRelease (0.018021)
13. thursdayRelease (0.016198)
14. topLeadProducer (0.015913)
15. Summer (0.014740)
16. saturdayRelease (0.014097)
17. Spring (0.011134)

Random Forest Regression

Why we chose this model.

- Ensemble method
 - It can perform both regression and classification tasks.
 - It can handle large datasets efficiently.
 - Provides a higher level of accuracy in predicting outcomes over the decision tree algorithm.
-

Extra Trees

Why we chose this model.

- Faster algorithm in terms of computation

Note:

- Random forest uses bootstrap replicas; it subsamples the input data with replacement, whereas Extra Trees use the whole original sample.
-

XGBoost Model

Why we chose this model.

- Scalable and highly accurate implementation of gradient boosting for boosted tree algorithms
- Built largely for energizing machine learning model performance and computational speed.

LightGBM

Why we chose this model.

- It's a distributed and efficient gradient boosting framework that uses tree-based learning.
- Histogram-based algorithm
- Places continuous values into discrete bins
 - Leads to faster training and more efficient memory usage.

Machine Learning Results

Random Tree Forest

Base Model:

Average Error: 1.5369
Accuracy = 89.478%

Model after Tuning:

Average Error: 1.3919
Accuracy = 90.394%
Improvement of 1.02%.

Extra Trees

Base Model:

Average Error: 1.6353
Accuracy = 88.852%

Model after Tuning:

Average Error: 1.4354
Accuracy = 90.139%
Improvement of 1.45%.

XGBoost Model

Base Model:

Average Error: 1.5313
Accuracy = 89.565%

Model after Tuning:

Average Error: 1.3809
Accuracy = 90.497%
Improvement of 1.041%.

LightGBM

Base Model:

Average Error: 1.4221
Accuracy = 90.214%

Model after Tuning:

Average Error: 1.3812
Accuracy = 90.481%
Improvement of 0.30%.

ML Model Conclusion

- XGBoost model performed the best after tuning.
- LightGBM model had the best “base” model of the four.
- The Machine Learning models implemented to predict revenue in this analysis were off by \$41M-\$43M.

Dashboard Presentation

[GitHub](#)

[Tableau](#)